

Hadoop Stock Exchange Case study

Mapreduce Hands-on

1] Write a map reduce program to split the NYSE_daily.txt into 4 files based on Ticker

```
hadoop jar nyse.jar casestudy1.SplitData /user/training/stock/input/NYSE_daily.txt /case1  
hadoop dfs -get /case1 output/
```

```
bitnami@ip-172-31-1-112:~$ hadoop jar nyse.jar casestudy1.SplitData /user/training/stock/input/NYSE_daily.txt /case1
16/07/16 13:17:41 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
16/07/16 13:17:42 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
16/07/16 13:17:42 INFO input.FileInputFormat: Total input paths to process : 1
16/07/16 13:17:42 INFO mapreduce.JobSubmitter: number of splits:1
16/07/16 13:17:42 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1468674155193_0001
16/07/16 13:17:43 INFO impl.YarnClientImpl: Submitted application application_1468674155193_0001
16/07/16 13:17:43 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1468674155193_0001/
16/07/16 13:17:43 INFO mapreduce.Job: Running job: job_1468674155193_0001
16/07/16 13:17:49 INFO mapreduce.Job: Job job_1468674155193_0001 running in uber mode : false
16/07/16 13:17:49 INFO mapreduce.Job: map 0% reduce 0%
16/07/16 13:17:54 INFO mapreduce.Job: map 100% reduce 0%
16/07/16 13:18:01 INFO mapreduce.Job: map 100% reduce 25%
16/07/16 13:18:04 INFO mapreduce.Job: map 100% reduce 50%
16/07/16 13:18:05 INFO mapreduce.Job: map 100% reduce 100%
16/07/16 13:18:06 INFO mapreduce.Job: Job job_1468674155193_0001 completed successfully
16/07/16 13:18:06 INFO mapreduce.Job: Counters: 50
  File System Counters
    FILE: Number of bytes read=3021950
    FILE: Number of bytes written=6634947
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=3194226
    HDFS: Number of bytes written=2907144
    HDFS: Number of read operations=15
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=8
  Job Counters
    Killed reduce tasks=1
    Launched map tasks=4
    Data-local map tasks=4
    Total time spent by all maps in occupied slots (ms)=2445
    Total time spent by all reduces in occupied slots (ms)=24253
    Total time spent by all map tasks (ms)=2445
    Total time spent by all reduce tasks (ms)=24253
    Total vcore-milliseconds taken by all map tasks=2445
    Total vcore-milliseconds taken by all reduce tasks=24253
    Total megabyte-milliseconds taken by all map tasks=2503680
    Total megabyte-milliseconds taken by all reduce tasks=24835072
```

2] Write a Map Reduce Program to find the maximum Variation for each stock ticker

```
hadoop jar nyse.jar casestudy2.MaxVariation /case1/p* /case2  
hadoop dfs -get /case2 output/
```

```
bitnami@ip-172-31-1-112:~$
bitnami@ip-172-31-1-112:~$ hadoop jar nyse.jar casestudy2.MaxVariation /case1/p* /case2
16/07/16 13:32:58 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
16/07/16 13:32:59 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
16/07/16 13:32:59 INFO input.FileInputFormat: Total input paths to process : 4
16/07/16 13:32:59 INFO mapreduce.JobSubmitter: number of splits:4
16/07/16 13:32:59 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1468674155193_0002
16/07/16 13:32:59 INFO impl.YarnClientImpl: Submitted application application_1468674155193_0002
16/07/16 13:32:59 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1468674155193_0002/
16/07/16 13:32:59 INFO mapreduce.Job: Running job: job_1468674155193_0002
16/07/16 13:33:05 INFO mapreduce.Job: Job job_1468674155193_0002 running in uber mode : false
16/07/16 13:33:05 INFO mapreduce.Job: map 0% reduce 0%
16/07/16 13:33:16 INFO mapreduce.Job: map 100% reduce 0%
16/07/16 13:33:21 INFO mapreduce.Job: map 100% reduce 100%
16/07/16 13:33:22 INFO mapreduce.Job: Job job_1468674155193_0002 completed successfully
16/07/16 13:33:22 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=1630270
    FILE: Number of bytes written=3851624
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=2907564
    HDFS: Number of bytes written=5979
    HDFS: Number of read operations=15
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=4
    Launched reduce tasks=1
    Data-local map tasks=4
    Total time spent by all maps in occupied slots (ms)=36174
    Total time spent by all reduces in occupied slots (ms)=2359
    Total time spent by all map tasks (ms)=36174
    Total time spent by all reduce tasks (ms)=2359
    Total vcore-milliseconds taken by all map tasks=36174
    Total vcore-milliseconds taken by all reduce tasks=2359
    Total megabyte-milliseconds taken by all map tasks=37042176
    Total megabyte-milliseconds taken by all reduce tasks=2415616
  Map-Reduce Framework
    Map input records=57391
    Map output records=57391
```